



# The GraphLab Technology Behind Create

Yucheng Low, PhD  
Chief Architect



# GraphLab Philosophy


## *Users-First* Architecture



# User



# Architecture




# Systems



User

Architecture

Systems



## Systems-First Architectures

Systems define constraints.  
Optimize for performance.




*PowerGraph*



User

Architecture

Systems



## Users-First Architectures

Users define constraints.  
Optimize for user interaction.



# What is a Users-First Architecture for Data Science?



# SFrame and SGraph

Built *by* data scientists,  
*for* data scientists.


Building on decades of database  
and systems research.



**SFrame**: Scalable Tabular  
Data Manipulation


		Title	Body
User			

**SGraph**: Scalable Graph  
Manipulation






Enabling users to **easily** and **efficiently** translate between both representations to get the best of both worlds.



**SFrame**: Scalable Tabular  
Data Manipulation

		Title	Body
User			


**SGraph**: Scalable Graph  
Manipulation



# **SFrame**: Scalable Tabular Data Manipulation

		Title	Body
User			

# **SGraph**: Scalable Graph Manipulation



# SFrame Design

## Pain Point #1: Resource Limits

Jobs fail because:

- Machine run out of memory
- Did not set Java Heap Size correctly
- Resource Configuration X needs to be bigger.



# SFrame Design

- **Graceful Degradation as 1<sup>st</sup> principle**
  - Always Works

## **Pain Point #2: Too Strict or Too Weak Schemas**

We want strong schema types.

We also want weak schema types.

Missing Values



# SFrame Design

- **Graceful Degradation as 1<sup>st</sup> principle**
  - Always Works
- **Rich Datatypes**
  - Strong schema types: int, double, string...
  - Weak schema types: list, dictionary



# SFrame Design

- **Graceful Degradation as 1<sup>st</sup> principle**
  - Always Works
- **Rich Datatypes**
  - Strong schema types: int, double, string...
  - Weak schema types: list, dictionary

## **Pain Point #3: Feature Manipulation**

Difficult or costly to inspect existing features and create new features.

Hard to perform data exploration.



# SFrame Design

- **Graceful Degradation as 1<sup>st</sup> principle**
  - Always Works
- **Rich Datatypes**
  - Strong schema types: int, double, string...
  - Weak schema types: list, dictionary
- **Columnar Architecture**
  - Easy feature engineering + Vectorized feature operations.
  - Immutable columns + Lazy evaluation
  - Statistics + visualization + sketches





# SFrame Python API Example

**Make a little SFrame of 1 column and 5 values:**

```
sf = gl.SFrame({'x':[1,2,3,4,5]})
```

**Normalizes the column x:**

```
sf['x'] = sf['x'] / sf['x'].sum()
```

**Uses a python lambda to create a new column:**

```
sf['x-squared'] = sf['x'].apply(lambda x: x*x)
```

**Create a new column using a vectorized operator:**

```
sf['x-cubed'] = sf['x-squared'] * sf['x']
```

**Create a new SFrame taking only 2 of the columns:**

```
sf2 = sf[['x','x-squared']]
```



# SFrame Querying

Supports most typical SQL SELECT operations using a Pythonic syntax.

## SQL


```
SELECT Book.title AS title, COUNT(*) AS authors
FROM Book
JOIN Book_author ON Book.isbn = Book_author.isbn
GROUP BY Book.title;
```

## SFrame Python

```
Book.join(Book_author, on='isbn')
    .groupby('title', {'authors':gl.aggregate.COUNT})
```



# SFrame Columnar Encoding





**Netflix Dataset,**  
**99M rows, 3 columns, ints**  
1.4GB raw  
289MB gzip compressed



# SFrame Columnar Encoding

SFrame File




## Type aware compression:

- Variable Bit length Encode
- Frame Of Reference Encode
- ZigZag Encode
- Delta / Delta ZigZag Encode
- Dictionary Encode
- General Purpose LZ4


**Netflix Dataset,**  
**99M rows, 3 columns, ints**  
1.4GB raw  
289MB gzip compressed



# SFrame Columnar Encoding



SFrame File



## Type aware compression:

- Variable Bit length Encode
- Frame Of Reference Encode
- ZigZag Encode
- Delta / Delta ZigZag Encode
- Dictionary Encode
- General Purpose LZ4


**Netflix Dataset,**  
**99M rows, 3 columns, ints**  
1.4GB raw  
289MB gzip compressed

User	→	176 MB	14.2 bits/int
Movie	→	<b>257 KB</b>	0.02 bits/int
Rating	→	47 MB	3.8 bits/int
-----			
Total	→	<b>223MB</b>	



# SFrame Columnar Encoding

SFrame File



## Type aware compression:

- Variable Bit length Encode
- Frame Of Reference Encode
- ZigZag Encode
- Delta / Delta ZigZag Encode
- Dictionary Encode
- General Purpose LZ4

**Netflix Dataset,**  
**99M rows, 3 columns, ints**  
1.4GB raw  
289MB gzip compressed

10s

User	→	176 MB	14.2 bits/int
Movie	→	257 KB	0.02 bits/int
Rating	→	47 MB	3.8 bits/int
-----			
Total	→	223MB	



# SFrames Distributed

**The choice of distributed or local execution is a question of query optimization.**


- Distributed Dataflow
- Columnar Query Optimizations
- Communicate columnar compressed blocks rather than row tuples.



# **SFrame**: Scalable Tabular Data Manipulation


		Title	Body
User			

# **SGraph**: Scalable Graph Manipulation





# SGraph: Scalable Graph Manipulation




# SGraph

- **SFrame backed** graph representation. Inherits SFrame properties.
  - Data types, External Memory, Columnar, compression, etc.
- Layout optimized for batch **external memory computation**.



# SGraph Layout

Vertex  
SFrames




Vertices Partitioned  
into  $p = 4$  SFrames.



# SGraph Layout

Vertices Partitioned  
into  $p = 4$  SFrames.

Vertex  
SFrames



<u>__id</u>	Name	Address	ZipCode
1011	John	...	98105
2131	Jack	...	98102



# SGraph Layout

Edges partitioned into  
 $p^2 = 16$  SFrames.

Vertex  
SFrames

Edge  
SFrames

1
2
3
4

(1,1)	(1,2)	(1,3)	(1,4)
(2,1)	(2,2)	(2,3)	(2,4)
(3,1)	(3,2)	(3,3)	(3,4)
(4,1)	(4,2)	(4,3)	(4,4)



# SGraph Layout

Edges partitioned into  
 $p^2 = 16$  SFrames.

Vertex  
SFrames

Edge  
SFrames

1
2
3
4

(1,1)	(1,2)	(1,3)	(1,4)
(2,1)	(2,2)	(2,3)	(2,4)
(3,1)	(3,2)	(3,3)	(3,4)
(4,1)	(4,2)	(4,3)	(4,4)



# SGraph Layout

Edges partitioned into  
 $p^2 = 16$  SFrames.

Vertex  
SFrames

Edge  
SFrames

1
2
3
4



(1,1)	(1,2)	(1,3)	(1,4)
(2,1)	(2,2)	(2,3)	(2,4)
(3,1)	(3,2)	(3,3)	(3,4)
(4,1)	(4,2)	(4,3)	(4,4)



# SGraph Layout

Vertex  
SFrames

Edge  
SFrames





# SGraph Layout

Vertex  
SFrames

Edge  
SFrames







# SGraph Layout

Vertex  
SFrames






Edge  
SFrames



# SGraph Layout

Vertex  
SFrames





Edge  
SFrames



# SGraph Layout

Vertex  
SFrames

Edge  
SFrames



# Deep Integration of SFrames and SGraphs

- Seamless interaction between graph data and table data.
- Queries can be performed easily across graph and tables.




Demo



**SFrame**: Scalable Tabular  
Data Manipulation

		Title	Body
User (ID)			

**SGraph**: Scalable Graph  
Manipulation



*User-first* architecture.

Built *by* data scientists,

*for* data scientists.

